

System Design 100 topics to learn

System Requirements:

1. Functional requirements:

These are requirements that define specific behavior or functions, often represented by use cases.

- **How to Define:** Start with user needs and business goals. Outline the primary tasks users should accomplish.
- **Working Backwards Approach:** Begin with the end goal and trace steps backwards to identify requirements.

2. High Availability:

Availability is a measure of system uptime.

- **Time-based:** E.g., 99.9% (or "three nines") uptime.
- **Count-based:** E.g., 1 failure allowed in 1 million requests.
- **Design Principles:** Redundancy, failover, replication, monitoring, and regular testing.
- **Processes:** Regular maintenance, backup procedures, disaster recovery planning.
- **SLO vs SLA:** Service Level Objective (SLO) is the target level of service. Service Level Agreement (SLA) is the contract with the customer promising certain SLOs.

3. Fault Tolerance, Resilience, Reliability:

- **Error:** An incorrect internal state.
- **Fault:** A component's incorrect behavior.
- **Failure:** When a system as a whole stops providing the required service.
- **Fault Tolerance:** Ability of a system to behave in a well-defined manner once a fault occurs.
- **Resilience:** Ability to return to a normal state after a failure.
- **Game Day vs Chaos Engineering:** Game Day is a planned simulation. Chaos Engineering is introducing random failures.
- **Reliability:** Probability a system will fail in a given period.

4. Scalability:

- **Vertical Scaling:** Increasing the resources of a single node.
- **Horizontal Scaling:** Adding more nodes.
- **Elasticity vs Scalability:** Elasticity is the ability to scale automatically based on demand.

5. Performance:

- **Latency:** Time taken to process a single operation.
- **Throughput:** Number of operations processed in a given time period.
- **Percentiles:** E.g., 95th percentile latency = under which 95% of the request latencies fall.
- **How to Increase:** Optimize algorithms, caching, distributed systems.

6. Durability:

Ensuring data is safely stored and retrievable.

- **Backup Types:**
 - **Full:** Everything is backed up.
 - **Differential:** Only changes since the last full backup.
 - **Incremental:** Only changes since the last backup of any type.
- **RAID:** Redundant Array of Independent Disks. A method for redundant storage.
- **Replication:** Copying data to ensure durability and availability.
- **Checksum:** A computed value to detect errors in data.

7. Consistency:

Ensuring all users see the same data.

- **Eventual Consistency:** Data will become consistent over time.
- **Linearizability:** Real-time consistency.
- **Monotonic Reads:** If a process reads the value of a data item, any successive reads by that process will always return that value or a more recent one.
- **Read-Your-Writes:** Guarantees that a write by a process is visible to a subsequent read by the same process.
- **Consistent Prefix Reads:** Guarantees that every read receives a prefix (some initial segment) of the writes, in order.

8. Maintainability, Security, Cost:

- **Maintainability:** Ease with which a system can be maintained. Includes monitoring, testing, and deployment strategies.
- **Security:**
 - **CIA Triad:** Confidentiality, Integrity, Availability.
 - **Identity and Permissions Management:** Who can do what.

- **Infrastructure Protection:** Protecting physical and virtual resources.
- **Data Protection:** Ensuring data confidentiality, integrity, and availability.
- **Cost Aspects:** Engineering, maintenance, hardware, and software costs.

9. Summary of System Requirements:

- **Most Popular Non-Functional Requirements:** Availability, Scalability, Performance, Security, Maintainability.

10. Regions, Availability Zones, etc.:

- **Regions:** Geographic areas that consist of multiple, isolated data centers.
- **Availability Zones:** Isolated locations within a region.
- **Data Centers:** Physical facilities housing servers.
- **Racks:** Structures holding multiple servers.
- **Servers:** Physical machines running software.

11. Physical Servers to Serverless:

- **Physical Servers:** Tangible hardware running software.
- **Virtual Machines (VMs):** Emulated computer systems.
- **Containers:** Lightweight VMs that share the same OS kernel and isolate the application processes from each other.
 - **Pros:** Lightweight, fast, consistent environment, scalable.
 - **Cons:** Less isolation than VMs.
- **Serverless:** Cloud-computing model which can run individual functions in response to events. Resources are fully managed by the cloud provider.
 - **Pros:** Scalability, no infrastructure management, cost-effective.
 - **Cons:** Cold starts, limited customization, statelessness.

12. Synchronous vs. Asynchronous Communication:

- **Synchronous:** Sender waits for the receiver to respond.
- **Asynchronous:** Sender sends the message and proceeds without waiting for a response.

13. Asynchronous Messaging Patterns:

- **Message Queuing:** Temporarily storing messages to be processed later.
- **Publish/Subscribe:** Senders (publishers) send messages to channels without specifying receivers. Receivers (subscribers) subscribe to channels they're interested in.

- **Competing Consumers:** Multiple consumers process messages from a single channel, distributing load.
- **Request/Response Messaging:** A two-way messaging pattern; a request is followed by a response.
- **Priority Queue:** Messages are processed based on priority rather than order of arrival.
- **Claim Check:** Large message is dropped off at a service and a claim check is provided to the requester, which can be used to get the full message later.

14. Network Protocols:

- **TCP (Transmission Control Protocol):** Reliable, connection-oriented protocol.
- **UDP (User Datagram Protocol):** Connectionless, no guarantee of message delivery.
- **HTTP (Hypertext Transfer Protocol):** Application protocol for distributed, collaborative, hypermedia systems.
 - **HTTP Request and Response:** Basic units of HTTP communication, where a client sends a request and a server responds.

15. Blocking vs. Non-blocking I/O:

- **Blocking I/O:** The process is blocked until data is available.
- **Non-blocking I/O:** The process continues even if the data is not available.
- **Thread per Connection Model:** Each connection gets a dedicated thread.
- **Thread per Request with Non-blocking I/O Model:** Threads are used only when processing a request.
- **Event Loop Model:** Single thread handles multiple connections using events.
- **Concurrency vs. Parallelism:** Concurrency is multiple tasks starting, running, and completing in overlapping time periods, while parallelism is multiple tasks running at exactly the same time.

16. Data Encoding Formats:

- **Textual Formats (e.g., JSON, XML):** Human-readable, more overhead.
- **Binary Formats (e.g., Protobuf, Avro):** Efficient, not human-readable.
- **Schema Sharing Options:** How schema (structure) of data is shared between sender and receiver.
- **Backward and Forward Compatibility:** Ensuring newer systems can read older data and older systems can read newer data, respectively.

17. Message Acknowledgment:

- **Safe Acknowledgment Modes:** Receiver confirms safe receipt of a message.

- **Unsafe Acknowledgment Modes:** Sender doesn't wait for a confirmation.

18. Deduplication Cache:

- **Local vs. External Cache:** Storage close to the CPU vs. external storage systems.
- **Adding Data to Cache:** Can be added explicitly by the application or implicitly by the caching system.
- **Cache Data Eviction:** Removing data from cache, can be based on size, time, or other criteria.
- **Expiration vs. Refresh:** Time after which data is removed vs. updating data periodically.

19. Metadata Cache:

- **Cache-aside Pattern:** Application is responsible for loading cache.
- **Read-through and Write-through Patterns:** Cache controls the data load and update.
- **Write-behind (Write-back) Pattern:** Data is written to cache and is asynchronously written back to the storage.

20. Queue:

- **Bounded and Unbounded Queues:** Limited vs unlimited size.
- **Circular Buffer (Ring Buffer):** Fixed-size data structure that uses a single, continuous region of memory.

21. Full and Empty Queue Problems:

- **Load Shedding:** Discarding excess traffic.
- **Rate Limiting:** Limiting the rate of requests by users.
- **Failed Requests Handling:** Strategies like retries or moving to dead-letter queues.
- **Backpressure:** Signal to the producer to slow down.
- **Elastic Scaling:** Dynamically adjusting resources based on demand.

22. Start with Something Simple:

- **Single Machine vs. Distributed System Concepts:** Many concepts in distributed systems (like caching, load balancing) can be understood with the analogy of single-machine concepts.
- **Interview Tip:** For system design interviews, start with a basic solution and then incrementally add complexity based on the requirements and constraints presented.

23. Blocking Queue and Producer-Consumer Pattern:

- **Producer-Consumer Pattern:** Separate components that produce and consume data, often operating at different rates.
- **Wait and Notify:** Mechanisms to synchronize between producers and consumers.

- **Semaphores:** A synchronization primitive used to control access to a shared resource.
- **Blocking Queue Applications:** Used in thread pooling, task scheduling, etc.

24. Thread Pool:

- **Pros:** Efficient use of resources, controlled resource usage, quick task start-up.
- **Cons:** Complexity, potential resource contention.
- **CPU-bound vs. I/O-bound tasks:** CPU-bound tasks spend most of their time using the CPU, while I/O-bound tasks spend time waiting for external operations (like disk or network operations).
- **Graceful Shutdown:** Properly terminating threads and releasing resources when shutting down a pool.

25. Big Compute Architecture:

- **Batch Computing Model:** Processing high volumes of data where a group of transactions is collected over a period.
- **Embarrassingly Parallel Problems:** Problems where little or no effort is needed to separate the problem into tasks to run in parallel.

26. Log:

- **Memory vs. Disk:** Storing logs in volatile memory (faster, but not persistent) vs. non-volatile disk storage (slower but persistent).
- **Log Segmentation:** Splitting logs into segments to manage them more efficiently.
- **Message Position (Offset):** The position of a message within a log.

27. Index:

- **How to Implement an Efficient Index for a Messaging System:** Use data structures like B-trees or Hash Indexes to allow for quick lookups and writes.

28. Time Series Data:

Storing and retrieving data points indexed in time order.

- **How to Store and Retrieve:** Use databases optimized for time series data, like InfluxDB, to ensure efficient storage and fast queries.

29. Simple Key-Value Database:

- **How to Build:** Utilize hash tables for quick lookups. Address issues like collision resolution.
- **Log Compaction:** Reducing the size of logs by removing redundant data.

30. B-tree Index:

- **Usage in Databases and Messaging Systems:** B-trees allow for efficient data storage and retrieval in databases and are used to index data in messaging systems.

31. **Embedded vs. Remote Database:**

- **Embedded Database:** Runs within the same address space as the application. E.g., SQLite.
- **Remote Database:** Separate from the application and is accessed over a network. E.g., PostgreSQL, MySQL.

32. **RocksDB:**

A high-performance embedded database for key-value data.

- **Memtable:** In-memory data structure where RocksDB writes are inserted.
- **Write-ahead Log (WAL):** A log that stores changes to data, ensuring durability.
- **Sorted Strings Table (SSTable):** Persistent, ordered immutable map of key-value pairs.

33. **LSM-tree vs. B-tree:**

- **Log-Structured Merge-Tree (LSM-tree):** Optimized for systems that write data more often than reading.
- **Write Amplification vs. Read Amplification:** Trade-offs between the cost of writing data vs. reading data.

34. **Page Cache:**

- **Increasing Disk Throughput:** Techniques like batching and zero-copy read can help in efficiently reading/writing to disks.

35. **Push vs. Pull:**

- **Push:** Server initiates sending data to the client.
- **Pull:** Client requests data from the server.

36. **Host Discovery:**

- **DNS:** The Domain Name System translates domain names to IP addresses.
- **Anycast:** Routing strategy where a single destination address has multiple routing paths to two or more endpoint destinations.

37. **Service Discovery:**

- **Server-side and Client-side Discovery Patterns:** Mechanisms where services find each other's network locations.
- **Service Registry:** A database containing the network locations of service instances.

38. **Peer Discovery:**

- **Peer Discovery Options:** How nodes in a network find each other.

- **Gossip Protocol:** Nodes randomly exchange information, which gradually propagates through the system.

39. Choosing a Network Protocol:

- **TCP, UDP, and HTTP:** Decision depends on use case requirements like reliability, speed, and connection state.

40. Video over HTTP:

- **Adaptive Streaming:** Adjusting video quality in real-time based on the viewer's network and playback conditions.

41. CDN (Content Delivery Network):

- **How to Use & Benefits:** CDNs distribute content across multiple servers to optimize access for users globally. They reduce latency and protect against traffic surges.
- **Point of Presence (POP):** Physical locations or access points that connect to the internet and host servers, acting as the local source for cached content.

42. Push and Pull Technologies:

- **Short Polling:** Client frequently asks the server for new data.
- **Long Polling:** Client requests information and waits for the server to respond.
- **Websocket:** Protocol that allows two-way communication with a server over a single, long-lived connection.
- **Server-Sent Events:** Server pushes updates to the client over an HTTP connection.

43. Push and Pull Technologies in Real-Life Systems:

- **Quiz:** Real-life examples would be chat applications (like Slack) using Websockets for real-time messaging, or stock trading apps using long polling for updating stock prices.

44. Large-scale Push Architectures:

- **C10K and C10M Problems:** Challenges of handling 10,000 and 10 million concurrent connections respectively.
- **Large-Scale Push Architectures:** Techniques like event-driven programming to handle large numbers of simultaneous connections.
- **Problems of Long-Lived Connections:** Resource management, detecting stale or "dead" connections, and handling dropped connections.

45. Building Reliable, Scalable, and Fast Systems:

- **Common Problems in Distributed Systems:** Network failures, machine failures, performance bottlenecks.
- **System Design Concepts:** Caching, load balancing, sharding, replication, and partitioning.

- **Three-Tier Architecture:** Presentation, logic, and data tiers.

46. Timeouts:

- **Fast Failures vs. Slow Failures:** Quickly failing a request vs. taking a long time before failing.
- **Connection and Request Timeouts:** Time limits set for establishing a connection or waiting for a response.

47. Handling Failed Requests:

- **Strategies:**
 - **Cancel:** Terminate the request.
 - **Retry:** Attempt the request again.
 - **Failover:** Redirect the request to another system or component.
 - **Fallback:** Use an alternate backup solution.

48. When to Retry:

- **Idempotency:** Ensuring that operations can be repeated without side effects.
- **AWS API Failures:** Not all API failures should be retried. For instance, validation errors shouldn't be retried.

49. How to Retry:

- **Exponential Backoff:** Increase the wait time between retries exponentially.
- **Jitter:** Adding randomness to retry intervals to spread out the load.

50. Message Delivery Guarantees:

- **At-most-once:** Messages are delivered once or not at all.
- **At-least-once:** Messages are delivered one or more times.
- **Exactly-once:** Messages are delivered precisely once.

51. Consumer Offsets:

- **Log-Based Messaging Systems:** Systems where messages are stored in order and consumers track their position using offsets.
- **Checkpointing:** Periodically saving the state of a process.

52. Batching:

- **Pros and Cons:** Batching can improve efficiency but may introduce latency.
- **Handling Batch Requests:** Ensuring that all items in a batch are processed, even if some fail.

53. Compression:

- **Pros and Cons:** Reduces data size but adds computational overhead.
- **Compression Algorithms:** Techniques like GZIP, Brotli, and LZ4, each with their trade-offs.

54. Scaling Message Consumption:

- **Single Consumer vs. Multiple Consumers:** One consumer can ensure order but is limited in throughput. Multiple consumers can increase throughput but introduce complexity in message ordering.
- **Problems with Multiple Consumers:** Ensuring order, handling duplicate processing, etc.

55. Partitioning in Real-Life Systems:

- **Pros and Cons:** Helps in horizontal scaling and distributing loads but can introduce complexities.
- **Applications:** Databases, messaging systems, etc.

56. Partitioning Strategies:

- **Lookup Strategy:** Direct lookup.
- **Range Strategy:** Based on a range of data values.
- **Hash Strategy:** Using a hash function for uniform distribution.

57. Request Routing:

- **Physical and Virtual Shards:** Actual data partitions vs logical partitions.
- **Routing Options:** Methods for directing a request to the appropriate server or data partition.

58. Rebalancing Partitions:

Ensuring data is evenly distributed across systems, especially when adding or removing nodes.

59. Consistent Hashing:

- **Implementation:** Uses a circular keyspace.
- **Advantages and Disadvantages:** Reduces the number of re-mapped keys when scaling but can lead to non-uniform data distribution.
- **Virtual Nodes:** Multiple hash values for a single node to ensure a more uniform distribution.

60. System Overload:

- **Importance:** Protecting against system overload ensures continued availability and prevents cascading failures across dependent systems.
- **Protection Mechanisms:** Load shedding, rate limiting, backpressure, and queuing.

61. Autoscaling:

- **Scaling Policies:**

- **Metric-based:** Scale based on metrics like CPU utilization or memory usage.
- **Schedule-based:** Scale based on known usage patterns or times.
- **Predictive:** Scale based on predicted future traffic patterns.
- **Autoscaling System Design:** Involves metrics collection, decision-making algorithms, and mechanisms to add or remove resources.

62. Load Shedding:

- **Implementation:** Prioritize and process only the most critical tasks, while dropping or delaying non-essential tasks.
- **Considerations:** Decide which requests to drop, how to notify clients, and when to trigger load shedding.

63. Rate Limiting:

- **Purpose:** To prevent any single user or client from overloading the system.
- **Implementation:** Token buckets, leaky buckets, and fixed window counters are popular methods.

64. Synchronous and Asynchronous Clients:

- **Admission Control Systems:** Mechanisms to decide whether to process a request immediately or defer it.
- **Blocking I/O vs. Non-Blocking I/O Clients:** Synchronous vs. Asynchronous client operations.

65. Circuit Breaker:

- **Finite-State Machine:** Mechanism where the circuit breaker maintains states (Closed, Open, Half-Open) to decide whether to process or deny requests.
- **Considerations:** Setting thresholds, reset intervals, and monitoring the health of downstream services.

66. Fail-Fast Design Principle:

- **Problems with Slow Services:** Slow services can cause chain reactions leading to larger system outages.
- **Solutions:** Implementing timeouts, circuit breakers, and other mechanisms to quickly fail and return control to the caller.

67. Bulkhead:

- **Implementation:** Isolating parts of a system so that if one fails, it doesn't bring down the whole system.
- **Applications:** Can be applied at various levels, from processes to servers to data centers.

68. Shuffle Sharding:

- **Implementation:** Combining sharding and redundancy to reduce the blast radius of failures.
- **Applications:** Reducing the impact of failures in multi-tenant systems.

69. Host Discovery:

- **DNS:** Translates domain names to IP addresses, enabling clients to discover hosts.
- **Anycast:** Uses DNS to route to the nearest or best-performing physical location.

70. Service Discovery:

- **Server-side and Client-side Discovery Patterns:** Mechanisms where services find each other's network locations, either with a central server or clients querying a service registry.
- **Service Registry:** Stores locations of service instances, allowing for service discovery.

71. Peer Discovery:

- **Peer Discovery Options:** Mechanisms like bootstrap lists or centralized directories.
- **Gossip Protocol:** Nodes share information about their peers, allowing for decentralized discovery.

72. Choosing a Network Protocol:

- **TCP vs. UDP vs. HTTP:** Depending on the requirements (reliability, connection setup, data size), one might be chosen over the other.

73. Network Protocols in Real-Life Systems:

Analyzing different scenarios and challenges to choose the best protocol. For example, streaming services might prefer UDP, while banking transactions would use TCP.

74. Video Over HTTP:

- **Adaptive Streaming:** The video quality is dynamically adjusted based on the viewer's network conditions.

75. CDN:

- **Benefits:** Faster content delivery by caching content closer to the end-users and reducing origin server load.
- **How It Works:** Content is distributed and stored across a network of servers. When a user requests content, it's served from the nearest cached server.

76. Push and Pull Technologies in Real-Life Systems:

Real-world applications like social media updates or live sports updates may use push technologies like WebSockets, while email clients might use pull technologies to check for new mails.

77. Large-Scale Push Architectures:

Design considerations for building systems that can handle millions of simultaneous connections.

78. **Timeouts:**

- **Importance:** Ensure that a system doesn't hang indefinitely.
- **Types:**
 - **Connection Timeouts:** Time a task will wait while attempting to establish a connection.
 - **Read/Write Timeouts:** Time a task will wait for a read or write operation to complete.
- **Handling:** Opt for graceful degradation of service, provide feedback to users, and use retries judiciously.

79. **Handling Failed Requests:**

- **Logging:** Essential for diagnostics. What failed, where, and why?
- **Retry Logic:** Implementing algorithms like exponential backoff.
- **Notifications:** Inform stakeholders or trigger systems about persistent failures.

80. **When to Retry:**

- **Transient vs Permanent Failures:** Temporary glitches (network blip) vs. consistent failures (invalid credentials).
- **Safety:** Ensuring retries don't exacerbate the problem.

81. **How to Retry:**

- **Immediate vs Delayed:** Waiting between retry attempts can be crucial.
- **Exponential Backoff:** Increasingly long waits between retries.
- **Jitter:** Introducing randomness to retry intervals.

82. **Message Delivery Guarantees:**

- **Importance:** Ensuring data isn't lost or seen multiple times.
- **Strategies:** Acknowledgments, transaction logs, checkpoints.

83. **Consumer Offsets:**

- **Checkpointing:** Storing position in a stream or log.
- **Benefits:** Allows consumers to pick up where they left off after failures.

84. **Batching:**

- **Benefits:** Can significantly reduce overheads and increase throughput.
- **Challenges:** Introducing latency, complexity in error handling.

85. Compression:

- **Trade-offs:** CPU time vs. bandwidth.
- **Use Cases:** Important in bandwidth-constrained environments or for storage savings.

86. Scaling Message Consumption:

- **Parallelism:** Distributing load across multiple consumers.
- **Ordering:** Ensuring that the sequence is maintained, if necessary.

87. Partitioning in Real-Life Systems:

- **Consistency vs Availability:** Deciding trade-offs based on system requirements.

88. Partitioning Strategies:

- **Consistency Hashing:** Minimizing reorganization of data when nodes are added or removed.

89. Request Routing:

- **Load Balancers:** Distributing incoming requests to prevent any single resource from being overwhelmed.
- **Content-based Routing:** Directing traffic based on the content of the request.

90. Rebalancing Partitions:

- **Importance:** Ensuring data and load are evenly distributed across nodes.
- **Strategies:** Manual interventions, algorithmic rebalancing.

91. Consistent Hashing:

- **Virtual Nodes:** A mechanism to ensure a more uniform distribution of data.

92. System Overload:

- **Monitoring and Alerts:** Real-time insights into system health.
- **Throttling:** Limiting user or service requests.

93. Autoscaling:

- **Cloud Services:** Many cloud providers offer autoscaling services that can automatically add or remove resources based on rules or machine learning.

94. Load Shedding:

- **Dynamic Adjustments:** Dropping non-critical tasks when the system is under heavy load.

95. Rate Limiting:

- **API Gateways:** Many come with built-in rate limiting features.

- **Client Feedback:** Informing clients when they're exceeding limits.

96. **Synchronous and Asynchronous Clients:**

- **Trade-offs:** Responsiveness vs. resource usage.
- **Use Cases:** Asynchronous operations for long-running tasks.

97. **Circuit Breaker:**

- **Resilience Patterns:** Preventing system failures from cascading.

98. **Fail-Fast Design Principle:**

- **Health Checks:** Regular checks to ensure services are operational.
- **Fallbacks:** Having backup systems or data sources in place.

99. **Bulkhead:**

- **Isolation:** Making sure failures in one part don't impact others.

100. **Shuffle Sharding:**

- **Resilience:** Improving system reliability by reducing the impact radius of failures.